

# EE 565: Computer Communication Networks I

## Lecture 7

Physical Layer 4 + Project 2

Winter Quarter 2022

# Today's Lecture: Physical Layer

- Physical Layer
  - Digital networking
  - Modulation
  - Characterization of Communication Channels
  - Fundamental Limits in Digital Transmission
  - Modems and Digital Modulation
  - Line Coding
  - Properties of Media and Digital Transmission Systems
  - Error Detection and Correction

# Today's Lecture: Physical Layer

- Physical Layer
  - Digital networking
  - Modulation
  - Characterization of Communication Channels
  - Fundamental Limits in Digital Transmission
  - Modems and Digital Modulation
  - Line Coding
  - Properties of Media and Digital Transmission Systems
  - Error Detection and Correction

# Error Control

- Channels introduce errors in digital communications
- Applications require certain reliability level
  - Data applications require error-free transfer
  - Voice & video applications tolerate some errors
- Error control may be needed to meet application requirement
- Error control ensures a data stream is transmitted to a certain level of accuracy despite errors
- Two basic approaches:
  - Error ***detection*** & retransmission (ARQ)
  - Forward error ***correction*** (FEC)

# Single Parity Code

- Information (7 bits): (0, 1, 0, 1, 1, 0, 0)
- Mini Quiz: Parity Bit?  
(“True”  $\rightarrow$  1, “False”  $\rightarrow$  0)  
 $b_8 = 0 + 1 + 0 + 1 + 1 + 0 = 1$   
Codeword (8 bits): (0, 1, 0, 1, 1, 0, 0, 1)
- If single error in bit 3? (0, 1, 1, 1, 1, 0, 0, 1)
  - # of 1's =5, odd  $\Rightarrow$  Error detected
- If errors in bits 3 and 5? (0, 1, 1, 1, 0, 0, 0, 1)
  - # of 1's =4, even  $\Rightarrow$  Error not detected

# Two-Dimensional Parity Check

- More parity bits to improve coverage
- Arrange information as columns
- Add single parity bit to each column
- Add a final “parity” column
- Used in early error control systems

1	0	0	1	0	0	0
0	1	0	0	0	0	1
1	0	0	1	0	0	0
1	1	0	1	1	1	0
1	0	0	1	1	1	1

Last column consists  
of check bits for each  
row

Bottom row consists of  
check bit for each column

# Error-detecting capability

1	0	0	1	0	0	0
0	0	0	0	0	0	1
1	0	0	1	0	0	0
1	1	0	1	1	1	0
1	0	0	1	1	1	1

One error

1	0	0	1	0	0	0
0	0	0	0	0	0	1
1	0	0	1	0	0	0
1	0	0	1	1	1	0
1	0	0	1	1	1	1

Two errors

1	0	0	1	0	0	0
0	0	0	1	0	0	1
1	0	0	1	0	0	0
1	0	0	1	1	1	0
1	0	0	1	1	1	1

Three errors

1	0	0	1	0	0	0
0	0	0	1	0	0	1
1	0	0	1	0	0	0
1	0	0	1	0	1	0
1	0	0	1	1	1	1

Four errors

1, 2, or 3 errors can always be detected; Not all patterns >4 errors can be detected

Arrows indicate failed check bits

# Checksum Calculation

The checksum  $\mathbf{b}_L$  is calculated as follows:

- Treating each 16-bit word as an integer, find

$$\mathbf{x} = \mathbf{b}_0 + \mathbf{b}_1 + \mathbf{b}_2 + \dots + \mathbf{b}_{L-1} \text{ modulo } 2^{16}-1$$

- The checksum is then given by:

$$\mathbf{b}_L = -\mathbf{x} \text{ modulo } 2^{16}-1$$

Thus, the headers must satisfy the following ***pattern*** at the receiver:

$$\mathbf{0} = \mathbf{b}_0 + \mathbf{b}_1 + \mathbf{b}_2 + \dots + \mathbf{b}_{L-1} + \mathbf{b}_L \text{ modulo } 2^{16}-1$$

- The checksum calculation is carried out in software using one's complement arithmetic



# Internet Checksum Example

## Use Modulo Arithmetic

- Assume 4-bit words
- Use mod  $2^4-1 (= 15)$  arithmetic
- $\underline{b}_0 = 1100 = 12$
- $\underline{b}_1 = 1010 = 10$
- $\underline{b}_0 + \underline{b}_1 = 12 + 10 = 7 \text{ mod } 15$
- $\underline{b}_2 = -7 = 8 \text{ mod } 15$
- Therefore
- $\underline{b}_2 = 1000$

## Use Binary Arithmetic

- Note  $16 = 1 \text{ mod } 15$
- So:  $10000 = 0001 \text{ mod } 15$
- leading bit wraps around

$$\begin{aligned} b_0 + b_1 &= 1100 + 1010 \\ &= 10110 \\ &= 10000 + 0110 \\ &= 0001 + 0110 \\ &= 0111 \\ &= 7 \end{aligned}$$

Take 1's complement

$$b_2 = -0111 = 1000$$

# Polynomial Codes

- Polynomials instead of vectors for codewords
- Polynomial arithmetic instead of check sums
- Implemented using shift-register circuits
- Also called *cyclic redundancy check (CRC)*
- Most data communications standards use polynomial codes for error detection
  - Have very simple hardware implementations
- Polynomial codes also basis for powerful error-correction methods

# Binary Polynomial Arithmetic

- Binary vectors map to polynomials

$$(i_{k-1}, i_{k-2}, \dots, i_2, i_1, i_0) \rightarrow i_{k-1}x^{k-1} + i_{k-2}x^{k-2} + \dots + i_2x^2 + i_1x^1 + i_0$$

Addition:

$$\begin{aligned}(x^7 + x^6 + 1) + (x^6 + x^5) &= x^7 + x^6 + x^6 + x^5 + 1 \\ &= x^7 + (1+1)x^6 + x^5 + 1 \\ &= x^7 + x^5 + 1 \quad \text{since } 1+1 = 0 \pmod{2}\end{aligned}$$

Multiplication:

$$\begin{aligned}(x+1)(x^2 + x + 1) &= x(x^2 + x + 1) + 1(x^2 + x + 1) \\ &= (x^3 + x^2 + x) + (x^2 + x + 1) \\ &= x^3 + 1\end{aligned}$$

# Binary Polynomial Division

- Division with Decimal Numbers

$$\begin{array}{r}
 34 \text{ ← quotient} \\
 35 \overline{) 1222} \text{ ← dividend} \\
 \underline{105} \phantom{00} \\
 172 \\
 \underline{140} \phantom{00} \\
 32 \text{ ← remainder}
 \end{array}$$

divisor

dividend = quotient x divisor + remainder

$$1222 = 34 \times 35 + 32$$

- Polynomial Division

$$\begin{array}{r}
 x^3 + x^2 + x \text{ ← divisor} \overline{) x^6 + x^5} \text{ ← dividend} \\
 \underline{x^6 + \phantom{x^5} x^4 + x^3} \\
 x^5 + x^4 + x^3 \\
 \underline{x^5 + \phantom{x^4} x^3 + x^2} \\
 x^4 + \phantom{x^3} x^2 \\
 \underline{x^4 + \phantom{x^3} x^2 + x} \\
 x \text{ ← remainder}
 \end{array}$$

$x^3 + x^2 + x = q(x)$  quotient

Note: Degree of  $r(x)$  is less than degree of divisor

$$x = r(x) \text{ remainder}$$

# Polynomial Coding

- k information bits define polynomial of degree k-1

$$i(x) = i_{k-1}x^{k-1} + i_{k-2}x^{k-2} + \dots + i_2x^2 + i_1x + i_0$$

$$g(x) = x^{n-k} + g_{n-k-1}x^{n-k-1} + \dots + g_2x^2 + g_1x + 1$$

- Find *remainder polynomial* of at most degree n-k-1

$$\begin{array}{r} q(x) \\ \hline g(x) \mid x^{n-k} i(x) \end{array} \qquad x^{n-k}i(x) = q(x)g(x) + r(x)$$

$r(x)$

- Define the *codeword polynomial* of degree n-1

$$\underbrace{b(x)}_{n \text{ bits}} = \underbrace{x^{n-k}i(x)}_{k \text{ bits}} + \underbrace{r(x)}_{n-k \text{ bits}}$$

Quiz Q: Find codeword if  $k=4$ ,  $n-k=3$

And: Generator polynomial:  $g(x) = x^3 + x + 1$

Information:  $(1,1,0,0)$        $i(x) = x^3 + x^2$

Encoding:  $x^3 i(x) = x^6 + x^5$

# Quiz Q: Find codeword if $k=4$ , $n-k=3$

And: Generator polynomial:  $g(x) = x^3 + x + 1$

Information:  $(1,1,0,0)$        $i(x) = x^3 + x^2$

Encoding:  $x^3 i(x) = x^6 + x^5$

$$\begin{array}{r} x^3 + x^2 + x \\ x^3 + x + 1 \overline{) x^6 + x^5} \\ \underline{x^6 + \phantom{x^5} x^4 + x^3} \phantom{+ x^2 + x} \\ x^5 + x^4 + x^3 \\ \underline{x^5 + \phantom{x^4} x^3 + x^2} \phantom{+ x} \\ x^4 + \phantom{x^3} x^2 \\ \underline{x^4 + \phantom{x^3} x^2 + x} \\ x \end{array}$$

# Quiz Q: Find codeword if $k=4$ , $n-k=3$

And: Generator polynomial:  $g(x) = x^3 + x + 1$

Information:  $(1,1,0,0)$        $i(x) = x^3 + x^2$

Encoding:  $x^3 i(x) = x^6 + x^5$

$$\begin{array}{r}
 x^3 + x^2 + x \\
 \hline
 x^3 + x + 1 \ ) \ x^6 + x^5 \\
 \underline{x^6 + \phantom{x^5} x^4 + x^3} \phantom{+ x^2 + x} \\
 x^5 + x^4 + x^3 \\
 \underline{x^5 + \phantom{x^4} x^3 + x^2} \phantom{+ x} \\
 x^4 + \phantom{x^3} x^2 \\
 \underline{x^4 + \phantom{x^3} x^2 + x} \\
 x
 \end{array}$$

$$\begin{array}{r}
 1110 \\
 \hline
 1011 \ ) \ 1100000 \\
 \underline{1011} \phantom{0000} \\
 1110 \\
 \underline{1011} \phantom{000} \\
 1010 \\
 \underline{1011} \phantom{00} \\
 010
 \end{array}$$

Transmitted codeword:

$$\begin{aligned}
 b(x) &= x^6 + x^5 + x \\
 \underline{b} &= (1,1,0,0,0,1,0)
 \end{aligned}$$




# The *Pattern* in Polynomial Coding

- All codewords satisfy the following **pattern**:

$$b(x) = x^{n-k}i(x) + r(x) = q(x)g(x) + r(x) + r(x) = q(x)g(x)$$

- All codewords are a multiple of  $g(x)$ !
- Receiver should divide received n-tuple by  $g(x)$  and check if remainder is zero
- If remainder is non-zero, then received n-tuple is not a codeword

# Undetectable error patterns



- $e(x)$  has 1's in error locations & 0's elsewhere
- Receiver divides the received polynomial  $R(x)$  by  $g(x)$
- Undetectable error: If  $e(x)$  is a multiple of  $g(x)$ , that is,  $e(x)$  is a non-zero codeword, then
- $R(x) = b(x) + e(x) = q(x)g(x) + q'(x)g(x)$
- *The set of undetectable error polynomials is the set of nonzero code polynomials*
- *Choose the generator polynomial so that selected error patterns can be detected.*

# Designing good polynomial codes

- Select generator polynomial so that likely error patterns are not multiples of  $g(x)$
- *Detecting Single Errors*
  - $e(x) = x^i$  for error in location  $i+1$
  - If  $g(x)$  has more than 1 term, it cannot divide  $x^i$
- *Detecting Double Errors*
  - $e(x) = x^i + x^j = x^i(x^{j-i} + 1)$  where  $j > i$
  - If  $g(x)$  has more than 1 term, it cannot divide  $x^i$
  - If  $g(x)$  is a *primitive* polynomial, it cannot divide  $x^m + 1$  for all  $m < 2^{n-k} - 1$  (Need to keep codeword length less than  $2^{n-k} - 1$ )
  - Primitive polynomials can be found by consulting coding theory books

# Standard Generator Polynomials

CRC = cyclic redundancy check

- CRC-8:  $= x^8 + x^2 + x + 1$  ATM
- CRC-16:  $= x^{16} + x^{15} + x^2 + 1$   
 $= (x + 1)(x^{15} + x + 1)$  Bisync
- CCITT-16:  $= x^{16} + x^{12} + x^5 + 1$  HDLC, XMODEM, V.41
- CCITT-32: IEEE 802, DoD, V.42  
 $= x^{32} + x^{26} + x^{23} + x^{22} + x^{16} + x^{12} + x^{11} + x^{10} + x^8 + x^7 + x^5 + x^4 + x^2 + x + 1$

# Hamming Codes

- Class of *error-correcting* codes
- Capable of **correcting** all *single-error* patterns
- Provably optimal for 1-bit errors
- Very less redundancy, e.g. 1-bit error proof – adds  $O(\log n)$  bits of redundancy for  $n$  bit sequences

## m=3      Hamming Code

- Information bits are  $b_1, b_2, b_3, b_4$
- Equations for parity checks  $b_5, b_6, b_7$

$$b_5 = b_1 \quad + b_3 + b_4$$

$$b_6 = b_1 + b_2 \quad + b_4$$

$$b_7 = \quad + b_2 + b_3 + b_4$$

- There are  $2^4=16$  codewords
- $(0,0,0,0,0,0,0)$  is a codeword

# My “simple” proof of optimality

Assume you got the following 7 bit sequences and make the following checks:

$$b_5 = b_1 \quad + b_3 + b_4$$

$$b_6 = b_1 + b_2 \quad + b_4$$

$$b_7 = \quad + b_2 + b_3 + b_4$$

Case	$b_5$ match	$b_6$ match	$b_7$ match
No error			
$b_1$ flipped			
$b_2$ flipped			
$b_3$ flipped			
$b_4$ flipped			
$b_5$ flipped			
$b_6$ flipped			
$b_7$ flipped			

# My “simple” proof of optimality


Assume you got the following 7 bit sequences and make the following checks:

$$\begin{aligned}b_5 &= b_1 + b_3 + b_4 \\b_6 &= b_1 + b_2 + b_4 \\b_7 &= b_1 + b_2 + b_3 + b_4\end{aligned}$$

Case	b <sub>5</sub> match	b <sub>6</sub> match	b <sub>7</sub> match
No error	✓	✓	✓
b <sub>1</sub> flipped	✗	✗	✓
b <sub>2</sub> flipped	✓	✗	✗
b <sub>3</sub> flipped	✗	✓	✗
b <sub>4</sub> flipped	✗	✗	✗
b <sub>5</sub> flipped	✗	✓	✓
b <sub>6</sub> flipped	✓	✗	✓
b <sub>7</sub> flipped	✓	✓	✗



# Why is Hamming a “good code”?




- Two valid bit sequences have a minimum distance of 3 bit flips
- Spheres of distance 1 around each codeword do not overlap
- If a single error occurs, the resulting n-tuple will be in a unique sphere around the original codeword
- Thus, receiver can correct erroneous reception back to original codeword

# Project 2

- Develop a UDP backend

# Develop a backend



- Hopefully, you have developed an operational frontend by now 😊
- Let's get to more complicated things...
- Develop a transport layer mechanism to implement a backend protocol
- Based on UDP
- Share content with friends  
(Similar to clearing each other's doubts on Canvas/Discord)

# What's new?

- First need an extra port for backend protocol

*java vodServer 8345 -> java vodServer 8345 8346*

- Support new URIs!

# Add Content URI

*/peer/add?path=<contentpath>&host=<IP/Hostname>&port=<backendport>&rate=<kbps>*

- This URI let the node knows that the specified remote content with the given path could be found on another node with the given IP address or hostname and port.

<http://localhost:8345/peer/add?path=content/video.ogg&host=pi.ece.uw.edu&port=8346&rate=1600>

- This URI tells the node that the file “content/video.ogg” could be found on pi.ece.uw.edu. The content path is relative to the directory where we start the peer. The back-end port is 8346. The average bit rate for the content is 1600 kilobits per second.

# View Content URI

`/peer/view/<contentpath>`

- This URI could be called by a web browser to retrieve the remote content using the specified path.

- After we enter this URIs (use the browser)

*`http://localhost:8345/peer/add?path=content/video.ogg&host=pi.ece.cmu.edu&port=8346`*

Then this URI should (immediately) play the video.

*`http://localhost:8345/peer/view/content/video.ogg`*

# Configuration URI

`/peer/config?rate=<bytes/s>`

- This URI specified the bandwidth limit (in bytes per second) that the backend transport can use for content transfer.

*`http://localhost:8345/peer/config?rate=8192`*

- Set a bandwidth limit on the backend transport to 8192 bytes/s.

# Status URI [Optional]

*/peer/status*

- This URI displays the average rate till now for the transfers on this backend port and IP address.

*http://localhost:8345/peer/status*

- Displays the average rate till now for the transfers on this backend port and IP address



# Requirements

- ***Multiplex***: You are allowed to use only one port for the backend transport (include sending & receiving). All backend traffic (even simultaneous transfers) should go through this port.
- ***Robustness***: Can reliably transfer the content, that is it can endure at least 5% packet loss
- ***Bandwidth Limit***: Can maintain the bandwidth specified in /peer/config option (can stay within 10% of the maximum bandwidth)
- ***Simultaneous transfer***: Can simultaneously transfer multiple (at least five) content requests at the same time.
- ***Content chunking***: When the same content is found on multiple peers (by sending multiple /peer/add requests with the same content path), the backend should be able to retrieve different parts of the content from different peers.

# Summary: How do I start?


- Start simple: UDP echoserver + client..
  - Transfer a file (you already know how to!)
- Create a header to send appropriate packets, multiplex, add sequence #, field for ACKs, etc.
- Implement flow control (window, ACKs)
- +Congestion control (updates to window)
- Note: Also do URI parsing, frontend-backend comm., etc..

# Outline

- What is the transport layer?
- Create/Destroy connection
- Flow control + Error recovery
- Congestion control

# Transport Protocols

- Lowest level end-to-end protocol (only TX, RX involved)
- Routers don't participate!



# What does it do?

- **Demultiplexing:** If there are 3 flows, how do I not mix them up? → Port #
- **Error detection:** If a packet is received incorrectly, how do I know? → checksum
- **Error recovery:** If an error happens, retransmit → ACKs+retransmit
- **Message boundaries** → Length field
- **In-order Delivery:** → Sequence #

Wait, I thought the MAC does this already?

→ *Only per-link, transport is end-to-end*

# What does it do?

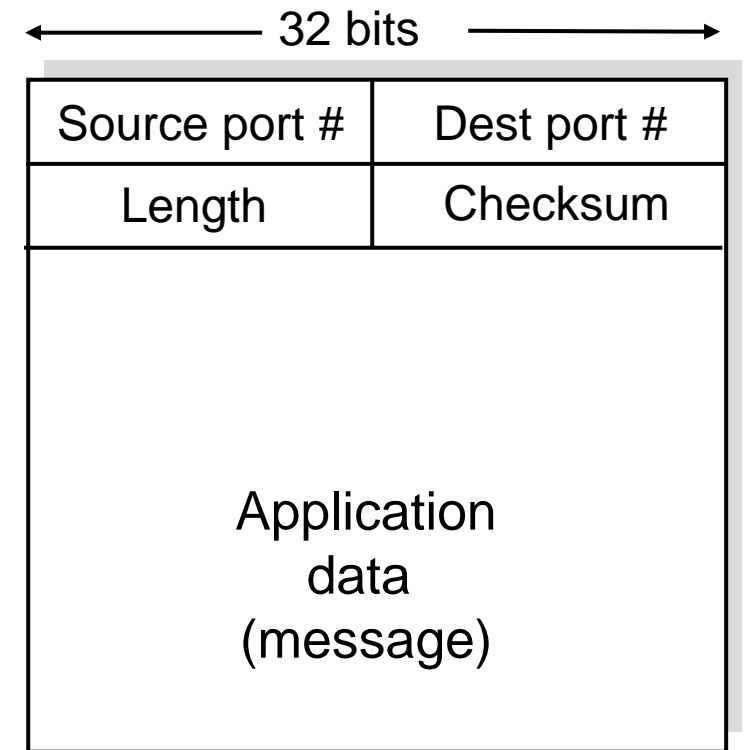
UDP

- **Demultiplexing:** If there are 3 flows, how do I not mix them up? → Port #
- **Error detection:** If a packet is received incorrectly, how do I know? → checksum
- **Error recovery:** If an error happens, retransmit → ACKs+retransmit
- **Message boundaries** → Length field
- **In-order Delivery:** → Sequence #

TCP

# UDP: User Datagram Protocol [RFC 768]

- “No frills,” “bare bones” Internet transport protocol
- Demultiplexing based on ports
- Optional checksum
  - One’s complement add (weak)



UDP segment format

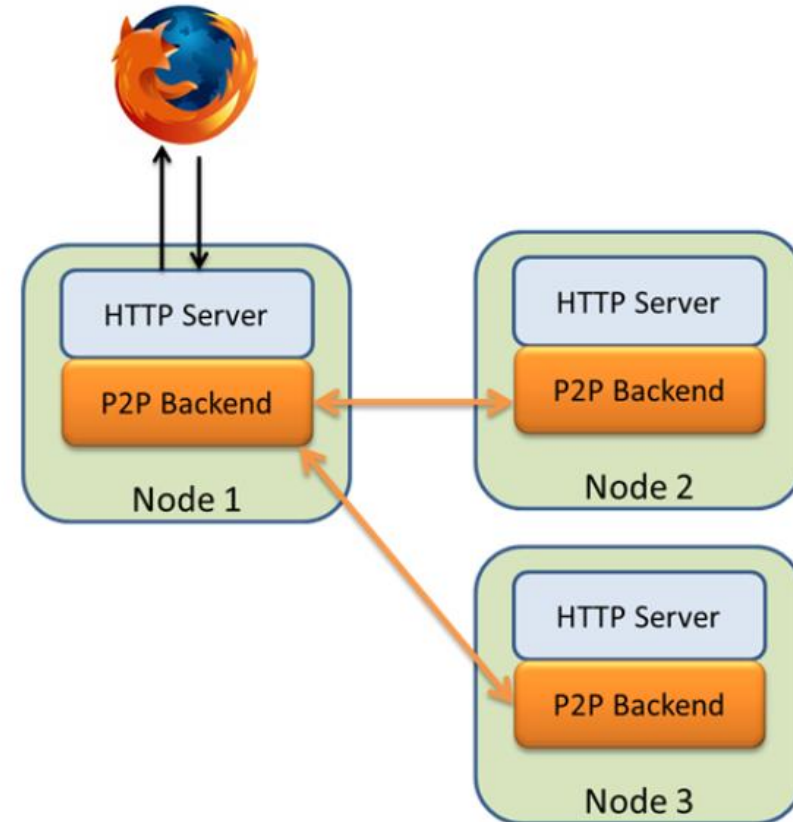
# UDP doesn't give a ...!

- Loss ... Not its problem?
- Reordering... Doesn't deal with it!
- Then what does it do?
  - Multiplexing + (Optional) Checksum
- Why is it useful?
  - Simple tasks (send 1 packet, e.g. beacon)
  - Building block for more complex protocols (e.g. project-2!)



# Project-2

- **Main task:** you need to develop a reliable transport protocol over UDP between backends (today)
- *Other stuff:* URI parsing, coordinating between HTTP server & backend (libraries)




# Ok.. What do I need to build?

- “Something like TCP over UDP.. (you can get away with less.. TCP-lite)”
  1. Create/Destroy connection
  2. Flow-control (in-order, error recovery)
  3. Congestion control

# 1. Create/Destroy Connection

Establishing 1-way Connection: Two-Way handshake

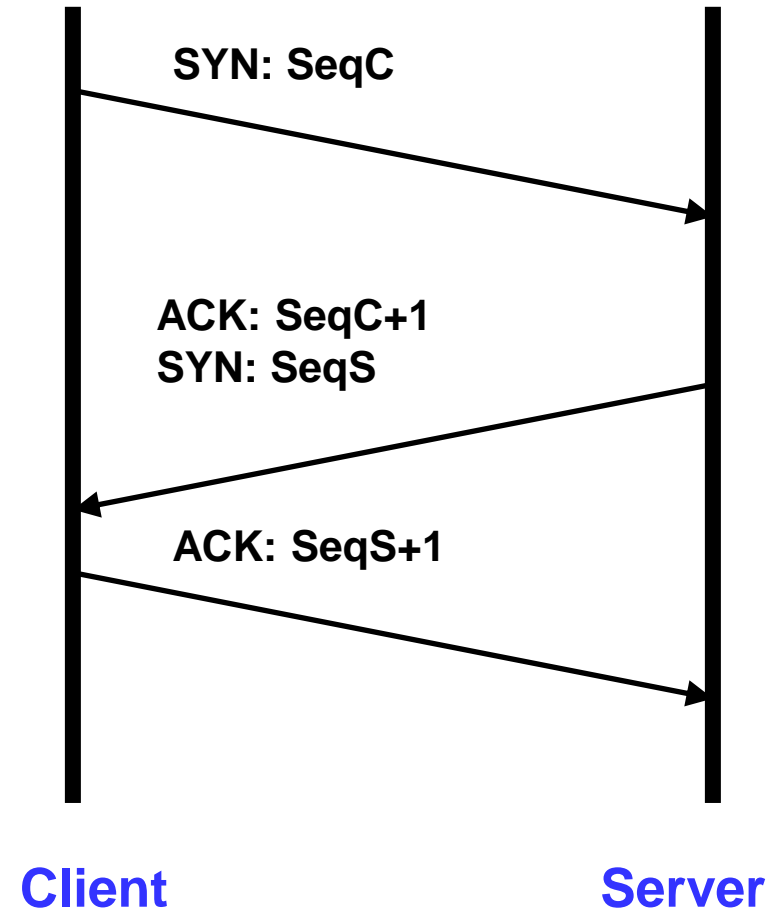
- Each side notifies other of starting sequence number it will use for sending
  - Why not simply chose 0?
    - Must avoid overlap with earlier incarnation
    - Security issues
- Each side acknowledges other's sequence number
  - SYN-ACK: Acknowledge sequence number + 1



# 1. Create/Destroy Connection

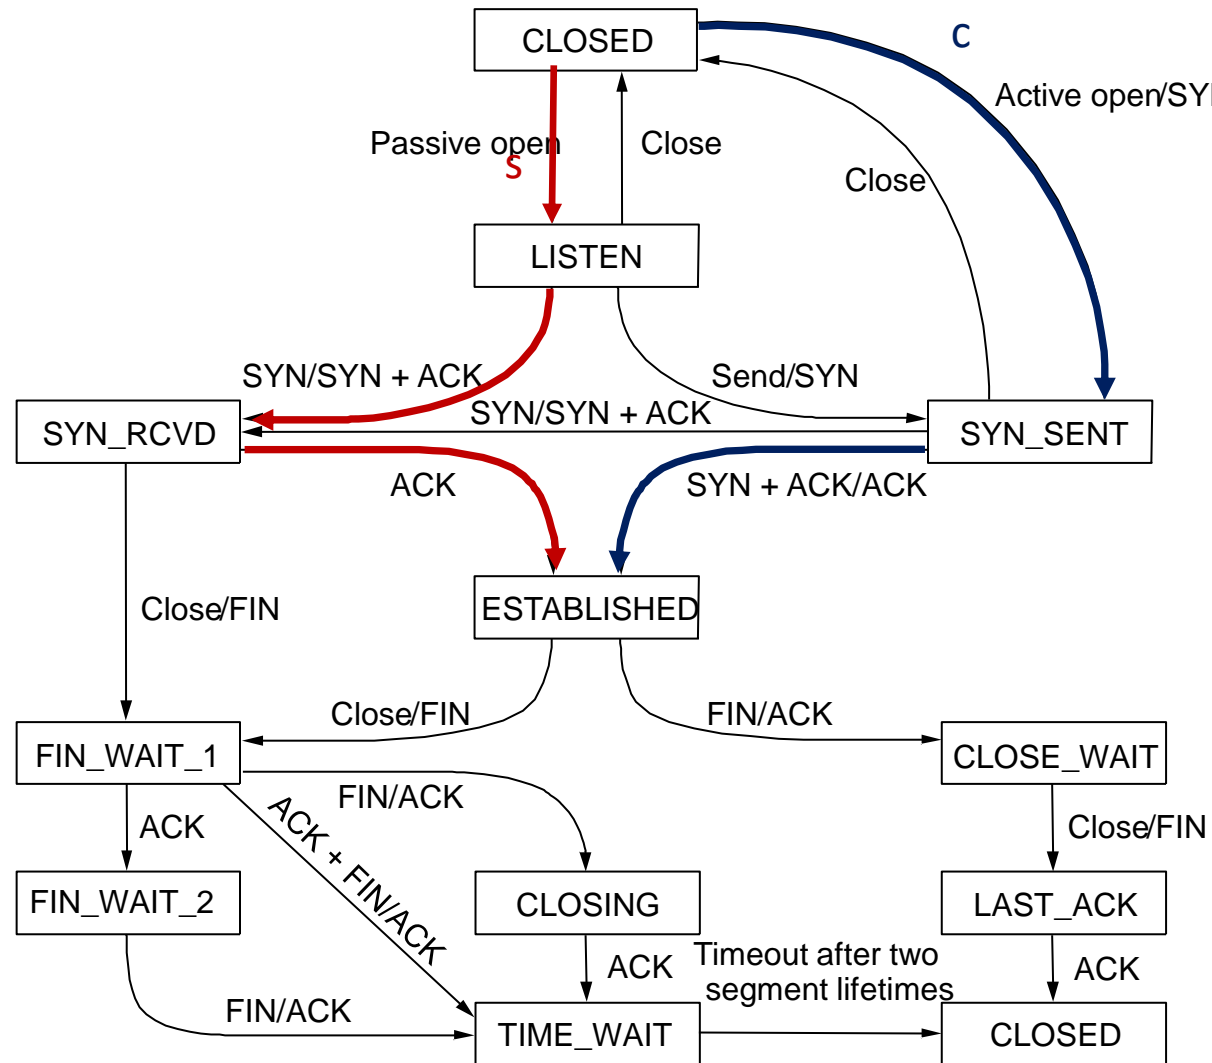
Establishing 2-way Connection: Three-Way handshake

- Optimize by combining second SYN with first ACK
- Lots of things can go wrong!
  - SYN lost
  - SYN-ACK lost
  - ACK lost
  - .. Time-out meaningfully



# 1. Create/Destroy Connection

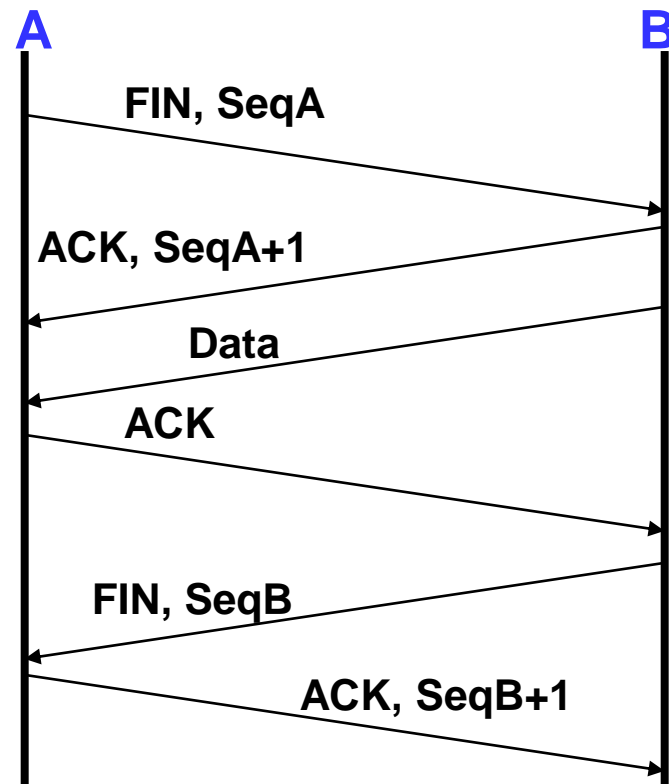
## TCP State Diagram: Connection Setup



# 1. Create/Destroy Connection

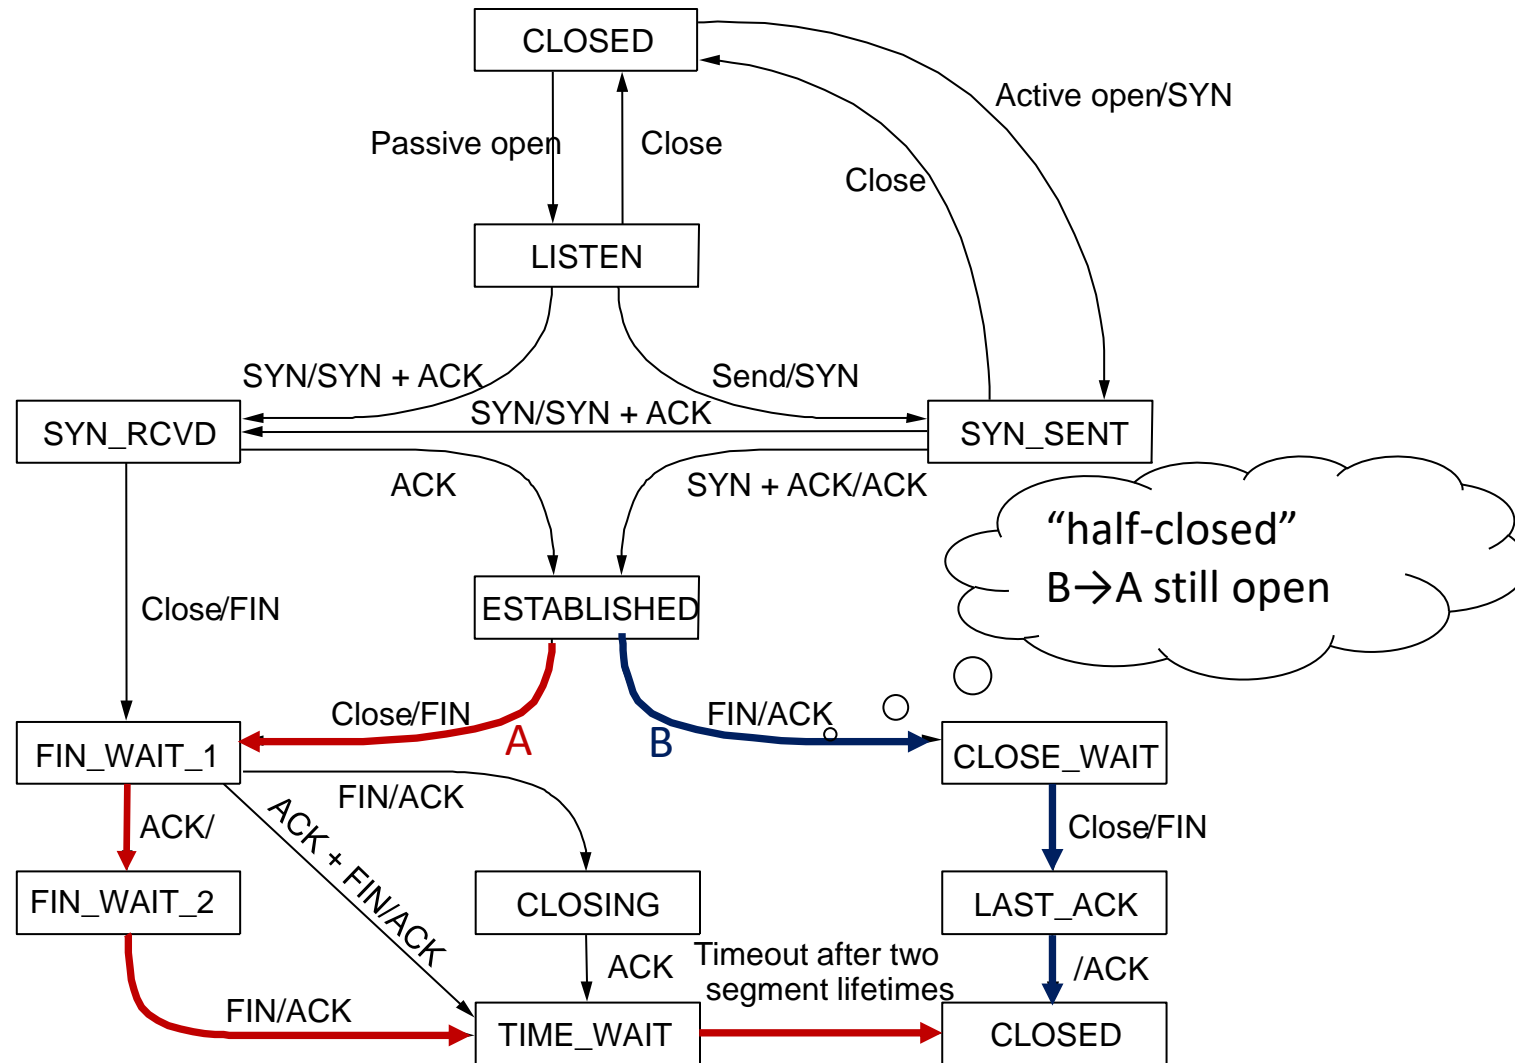
## Tearing Down Connection

- Either side can initiate tear down
  - Send FIN signal
  - “I’m not going to send any more data”
- Other side can continue sending data
  - Half open connection
  - Must continue to acknowledge
- Acknowledging FIN
  - Acknowledge last sequence number + 1



# 1. Create/Destroy Connection

## TCP State Diagram: Tearing Down Connection



# Ok.. What do I need to build?

- “Something like TCP over UDP.. (you can get away with less.. TCP-lite)”

- ✓ 1. Create/Destroy connection
2. Flow-control (in-order, error recovery)
3. Congestion control



# Ok.. What do I need to build?

- “Something like TCP over UDP.. (you can get away with less.. TCP-lite)”

- ✓ 1. Create/Destroy connection
2. Flow-control (in-order, error recovery)
3. Congestion control


# Goals here..

- Ensure in-order delivery...  
→ Sequence numbers
- Ensure you know packets are lost  
→ ACKs
- Recover from loss → Retransmit

# Example: TCP Header


Source port		Destination port	
Sequence number			
Acknowledgement			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			

# ACKs are important




# Review: Stop and Wait is BAD

- Simplest ARQ protocol
- Send a packet, stop and wait until ACK arrives
- Inefficient!



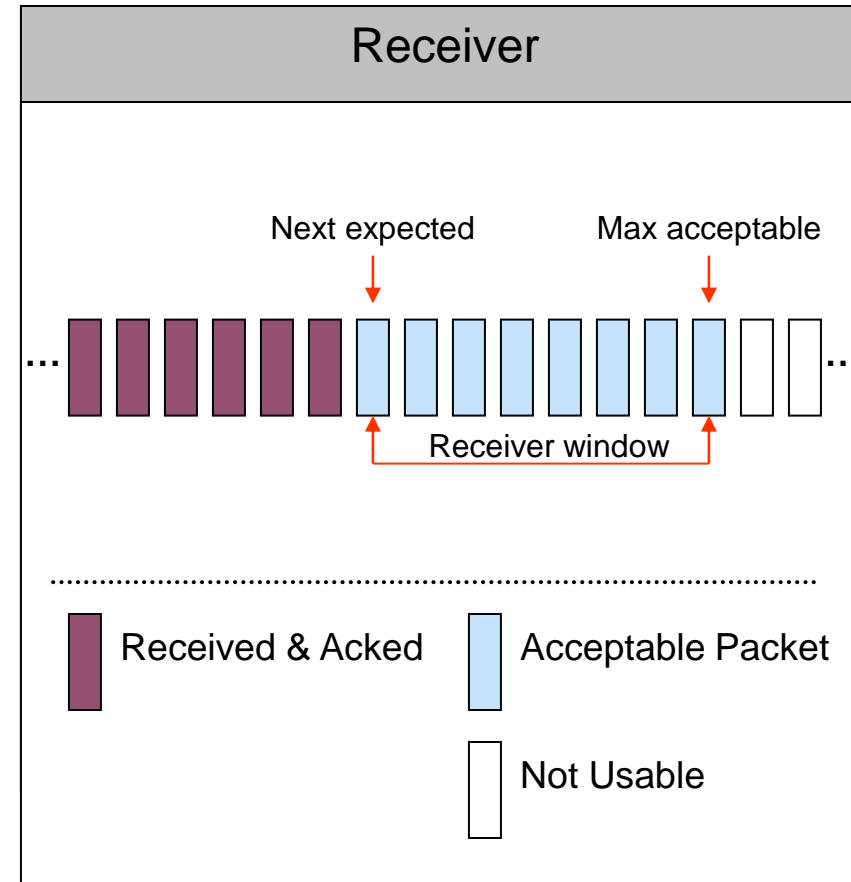
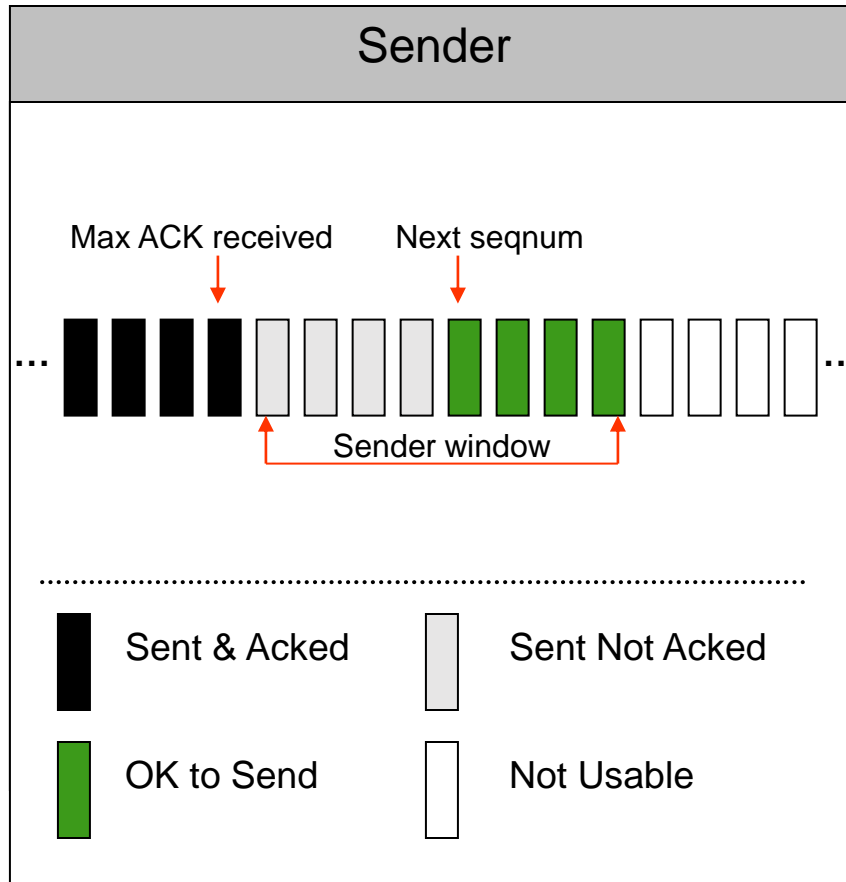
Correct approach:  
Window = Bandwidth-Delay Product



$$\text{Max Throughput} = \frac{\text{Window Size}}{\text{Roundtrip Time}}$$

# Sliding Window

## Sender/Receiver State



# Sliding Window Algorithm 101

- @TX: If new ACK is received (in order)
  - Increment MAX ack received
  - Send next packet
- @RX: If new pkt is received (in order)
  - Increment next expected
  - Increment MAX acceptable
- What about loss / out-of-order?
  - Many options! (up to you)
  - Simplest: Both TX & RX share their queue state



# Important Detail: Timeout = ?

- Wait at least one RTT before retransmitting
- Importance of accurate RTT estimators:
  - Low RTT estimate
    - unneeded retransmissions
  - High RTT estimate
    - poor throughput
- RTT estimator must adapt to change in RTT
- How to pick timeout?
  - Up to you!
  - First cut: Estimate based on history.. Be conservative

# Ok.. What do I need to build?

- “Something like TCP over UDP.. (you can get away with less.. TCP-lite)”

- ✓ 1. Create/Destroy connection
- ✓ 2. Flow-control (in-order, error recovery)
3. Congestion control

# Congestion Control

How do I pick window?

- Too small → Underutilization
- Too big → Congestion!
- “Just right” ... but how do we find it?

# Idea: Additive Increase Multiplicative Decrease (AIMD)

- Initialize Window
- Every RTT...
  - If all packets received:  $W = W + 1$
  - If any packet dropped:  $W = W/2$
- Why not: MIAD or AIAD or MIMD?

# Proof by Example

- Assume two users  $W_0=1$ ,  $W_1=5$ ... max 10 packets in n.w.

- AIAD:

$(1,5) \rightarrow (2,6) \rightarrow (3,7) \rightarrow (4,8) \rightarrow (3,7) \rightarrow (4,8) \rightarrow \dots$ (repeat)

- MIAD:

$(1,5) \rightarrow (2,10) \rightarrow (1,9) \rightarrow (2,18) \rightarrow (0,19) \rightarrow \dots$ (congest!!)

- MIMD:


$(1,5) \rightarrow (2,10) \rightarrow (1,5) \rightarrow (2,10) \rightarrow \dots$ (repeat)

- AIMD:

$(1,5) \rightarrow (2,6) \rightarrow (3,7) \rightarrow (4,8) \rightarrow (2,4) \rightarrow (3,5) \rightarrow (4,6) \rightarrow (5,7) \rightarrow$   
 $(3,4) \rightarrow (4,5) \rightarrow (5,6) \rightarrow (3,3) \rightarrow (4,4) \rightarrow (5,5) \rightarrow (6,6) \rightarrow \text{repeat}$

# TCP Congestion Control: Implicit Feedback and AIMD

- Distributed, fair and efficient
- Packet loss is seen as sign of congestion and results in a multiplicative rate decrease: factor of 2
- TCP periodically probes for available bandwidth by increasing its rate: by one packet per RTT



# Implementing in practice

- Per packet..
  - $W = W+1/W...$  (so  $W=W+1$  in one RTT)
  - $W = W/2$  whenever a timeout occurs
- Many optimizations possible
  - Back-off if you detect loss from feedback
  - Increase  $W$  faster initially, e.g.  $W=2W$  in one RTT
  - Many more... (up to you!)

# Ok.. What do I need to build?

- “Something like TCP over UDP.. (you can get away with less.. TCP-lite)”

- ✓ 1. Create/Destroy connection
- ✓ 2. Flow-control (in-order, error recovery)
- ✓ 3. Congestion control



# Summary: How do I start?

- Start simple: UDP echoserver + client..
  - Transfer a file (you already know how to!)
- Create a header to send appropriate packets, multiplex, add sequence #, field for ACKs, etc.
- Implement flow control (window, ACKs)
- +Congestion control (updates to window)
- Note: Also do URI parsing, frontend-backend comm., etc..